


Checkmate with AI

Alexandros Stathakopoulos, Mohanad Kandil
Technische Universität München
Heilbronn, 29. Januar 2025



TUM Uhrenturm

Agenda

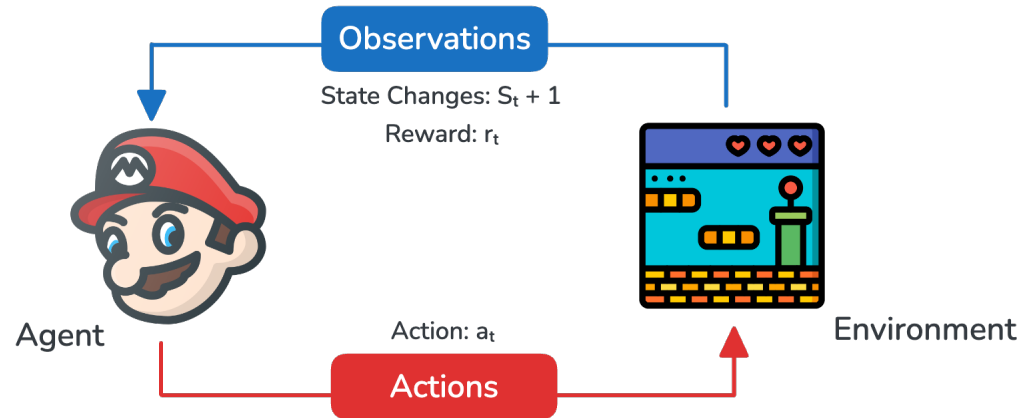
1. Introduction
2. Model Overview
3. Algorithm Explanation
4. Code Structure & Implementation
5. Algorithm Performance & Results
6. Implications & Limitations
7. Conclusion & Q&A

Introduction - Background & Motivation

- Problem Statement: What are we solving?
- Why is reinforcement learning important?
- Brief mention of prior research
- Project Goal: Train an AI agent using RL

Model Overview

- What is our RL model trying to achieve?
- Type of RL Model: Q-Learning & Deep Q-Network (DQN)
- High-Level Concept: Interaction between Agent and Environment



Algorithm Explanation - Deep Q-Learning

- Reinforcement learning technique to find optimal action-value function $Q(s, a)$
- Uses a neural network instead of a Q-table
- Action Selection: ϵ -greedy policy for exploration vs exploitation
- Experience Replay to improve stability

Code Structure & Implementation

- Overview of Repository Structure
- Key Libraries: Python, TensorFlow/PyTorch, NumPy, OpenAI Gym
- Training and Evaluation Pipeline

Algorithm Performance & Results

- Performance Metrics: Reward over Episodes
- Training Loss over Iterations
- Baseline Model Comparison

Implications

- Application in Finance, Robotics, Gaming, and Cloud Optimization
- Enhancements for real-world decision-making
- Future research directions: Combining RL with Transformer models

Limitations

- High computational cost
- Sample inefficiency
- Lack of interpretability
- Generalization issues across different environments

Conclusion & Summary


- Key Takeaways: Insights from Model Performance
- Challenges & Future Work

Conclusion & Summary

- Key Takeaways: Insights from Model Performance
- Challenges & Future Work

Any Questions?

Thank You



THE REASON I AM SO INEFFICIENT

xkcd 1445